

File `countrypop.csv` contains total population numbers (in thousands) for 40 different countries for the years 1950, 1960, 1970, 1980, 1990, 2000, and 2010. Build and compare two different varying-coefficient hierarchical normal regression models for the log-scale numbers, using JAGS and rjags.

```
data<- read.csv("./countrypop.csv", header=TRUE)
head(data,10)
```


##	Country	Year1950	Year1960	Year1970	Year1980	Year1990	Year2000
## 1	Andorra	6.198	13.410	24.275	36.063	54.508	65.390
## 2	Aruba	38.052	54.208	59.070	60.097	62.152	90.866
## 3	Austria	6936.442	7070.773	7516.238	7609.750	7723.954	8069.276
## 4	Azerbaijan	2927.926	3895.398	5180.032	6150.735	7242.758	8122.743
## 5	Bahrain	115.612	162.429	212.607	359.897	495.927	664.610
## 6	Bangladesh	37894.671	48013.505	64232.486	79639.498	103171.957	127657.862
## 7	Belarus	7745.004	8124.881	8913.549	9569.847	10151.135	9871.635
## 8	Cameroon	4307.021	5176.920	6519.754	8621.409	11780.086	15513.944
## 9	Comoros	159.459	191.122	230.055	307.831	411.598	542.358
## 10	Croatia	3850.294	4192.641	4423.069	4598.125	4776.370	4428.075
##	Year2010						
## 1	84.454						
## 2	101.665						
## 3	8409.945						
## 4	9032.465						
## 5	1240.864						
## 6	147575.433						
## 7	9420.576						
## 8	20341.236						
## 9	689.696						
## 10	4328.163						

(a)

(i) On the same set of axes, plot segmented lines, one for each country, representing the population (in thousands) versus the year (1950, 1960, ...). Distinguish the lines for different countries by using different colors or line types.

```
xs <- seq(1950, 2010, by=10)
plot(xs, data[1, -1], type="l", col=rainbow(nrow(data))[1], ylim=c(0, max(data[-1])), xlab="Year", ylab="Population")
for (i in 2:nrow(data)) {
  lines(xs, data[i, -1], col=rainbow(nrow(data))[i])
}
```


## Population Trends 1950–2010



(ii) Repeat the previous part using the natural logarithm of the population numbers.

```
data[, -1] <- log(data[, -1])
plot(xs, data[1, -1], type="l", col=rainbow(nrow(data))[1], ylim=c(0, max(data[-1])), xlab="Year", ylab="Log Population")
for (i in 2:nrow(data)) {
  lines(xs, data[i, -1], col=rainbow(nrow(data))[i])
}
```

## Log Population Trends 1950–2010



Let  $y_{ij}$  be the natural logarithm of the population (in thousands) of country  $j$  in the year indexed by  $i$  ( $i = 1, \dots, 7$  corresponding to 1950, ..., 2010), and  $j = 1, \dots, 40$ . For each country, let the log-population be modeled as a simple linear regression on the centered year index:

$$y_{ij} | \beta^{(j)}, \sigma_y^2 \sim \text{indep. } N(\beta_1^{(j)} + \beta_2^{(j)}(x_i - \bar{x}), \sigma_y^2)$$

where

$$\beta^{(j)} = \begin{pmatrix} \beta_1^{(j)} \\ \beta_2^{(j)} \end{pmatrix}, \quad j = 1, \dots, 40$$

and

$$x_i = i, \quad i = 1, \dots, 7$$

Note that the coefficients are allowed to depend on the country, but the variance is not.


(b) Let  $\hat{\beta}_1^{(j)}$  and  $\hat{\beta}_2^{(j)}$  be the ordinary least squares estimates of  $\beta_1^{(j)}$  and  $\beta_2^{(j)}$ , estimated for country  $j$ .

(i) Produce a scatterplot of the pairs  $(\hat{\beta}_1^{(j)}, \hat{\beta}_2^{(j)})$ ,  $j = 1, \dots, 40$ .

```
years <- c(1950, 1960, 1970, 1980, 1990, 2000, 2010)
centered_years <- years - 1980
betahat <- matrix(NA, nrow = nrow(data), ncol = 2)
for (j in 1:nrow(data)) {
  population_data <- t(data[j, -1])
  fit <- lsfit(x = centered_years, y = population_data)
  betahat[j, ] <- fit$coef
}
```

```
plot(betahat[, 1], betahat[, 2], xlab = "betahat1", ylab = "betahat2", main = "Scatterplot of Beta Coef")
```

## Scatterplot of Beta Coefficients



(ii) Give the average (sample mean) of  $\hat{\beta}_1^{(j)}$  and also of  $\hat{\beta}_2^{(j)}$ .

```
means = apply(betahat, 2, mean)
cat("mean_betahat1: ", means[1], "\n", "mean_betahat2: ", means[2])
```

```
## mean_betahat1: 8.159197
## mean_betahat2: 0.01991429
```

(iii) Give the sample variance of  $\hat{\beta}_1^{(j)}$  and also of  $\hat{\beta}_2^{(j)}$ .

```
var_betahat1 = var(betahat)[1,1]
var_betahat2 = var(betahat)[2,2]
cat("var_betahat1: ", var_betahat1, "\n", "var_betahat2: ", var_betahat2, "\n")
```

```
## var_betahat1: 4.946005
## var_betahat2: 0.0002313094
```

(iv) Give the sample correlation between  $\hat{\beta}_1^{(j)}$  and  $\hat{\beta}_2^{(j)}$ .

```
betahat_cov <- cov(betahat[, 1], betahat[, 2])
cat("betahat_cov: ", betahat_cov, "\n")
```

```
## betahat_cov: -0.00566316
```

(c) Consider the bivariate prior

$$\beta^{(j)} | \mu_\beta, \Sigma_\beta \sim \text{iid } N(\mu_\beta, \Sigma_\beta)$$

with

$$\mu_{\beta} = \begin{pmatrix} \mu_{\beta 1} \\ \mu_{\beta 2} \end{pmatrix}$$

and covariance matrix

$$\Sigma_{\beta} = \begin{pmatrix} \sigma_{\beta 1}^2 & \rho \sigma_{\beta 1} \sigma_{\beta 2} \\ \rho \sigma_{\beta 1} \sigma_{\beta 2} & \sigma_{\beta 2}^2 \end{pmatrix}$$

with hyperpriors

$$\mu_{\beta} \sim N(0, 1000^2 I)$$

and

$$\Sigma_{\beta}^{-1} \sim \text{Wishart}_2((\Sigma_0^{-1})/2)$$

in the notation used in the lecture videos. For your analysis, use

$$\Sigma_0 = \begin{pmatrix} 5 & 0 \\ 0 & 0.02 \end{pmatrix}$$

based on preliminary analyses. Let the prior on  $\sigma_y^2$  be

$$\sigma_y^2 \sim \text{Inv-gamma}(0.0001, 0.0001)$$

(i) List an appropriate JAGS model. Make sure to create nodes for  $\Sigma_{\beta}^{-1}$ ,  $\rho$ , and  $\sigma_y^2$ .

```
d1 <- list(population = data[, -1],
  year = xs,
  mubeta0 = c(0, 0),
  Sigmamubetainv = rbind(c(0.000001, 0),
    c(0, 0.000001)),
  Sigma0 = rbind(c(5, 0),
    c(0, 0.02)))

inits1 <- list(list(sigmasqyinv = 10, mubeta = c(1000, 1000),
  Sigmabetainv = rbind(c(100, 0),
    c(0, 100))),
  list(sigmasqyinv = 0.001, mubeta = c(-1000, 1000),
    Sigmabetainv = rbind(c(100, 0),
      c(0, 100))),
  list(sigmasqyinv = 10, mubeta = c(1000, -1000),
    Sigmabetainv = rbind(c(0.001, 0),
      c(0, 0.001))),
  list(sigmasqyinv = 0.001, mubeta = c(-1000, -1000),
    Sigmabetainv = rbind(c(0.001, 0),
      c(0, 0.001))))

library(rjags)

## Loading required package: coda
## Linked to JAGS 4.3.2
## Loaded modules: basemod,bugs

#print the model
cat (readLines('countrypop_1.bug'), sep= '\n')
```

```
## data {
##   dimY <- dim(population)
##   yearcent <- year - mean(year)
## }
##
## model {
##   for (j in 1:dimY[1]) {
##     for (i in 1:dimY[2]) {
##       population[j,i] ~ dnorm(beta[1,j] + beta[2,j]*yearcent[i], sigmasqyinv)
##     }
##     beta[1:2,j] ~ dmnorm(mubeta, Sigmabetainv)
##   }
##   mubeta ~ dmnorm(mubeta0, Sigmamubetainv)
##   Sigmabetainv ~ dwish(2*Sigma0, 2)
##   sigmasqyinv ~ dgamma(0.0001, 0.0001)
##   Sigmabeta <- inverse(Sigmabetainv)
##   rho <- Sigmabeta[1,2] / sqrt(Sigmabeta[1,1] * Sigmabeta[2,2])
##   sigmasqy <- 1/sigmasqyinv
## }
```

```
m1 <- jags.model("countrypop_1.bug", d1, inits1, n.chains=4, n.adapt=5000)
```

```
## Compiling data graph
##   Resolving undeclared variables
##   Allocating nodes
##   Initializing
##   Reading data back into data table
## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 280
##   Unobserved stochastic nodes: 43
##   Total graph size: 1004
##
## Initializing model
```

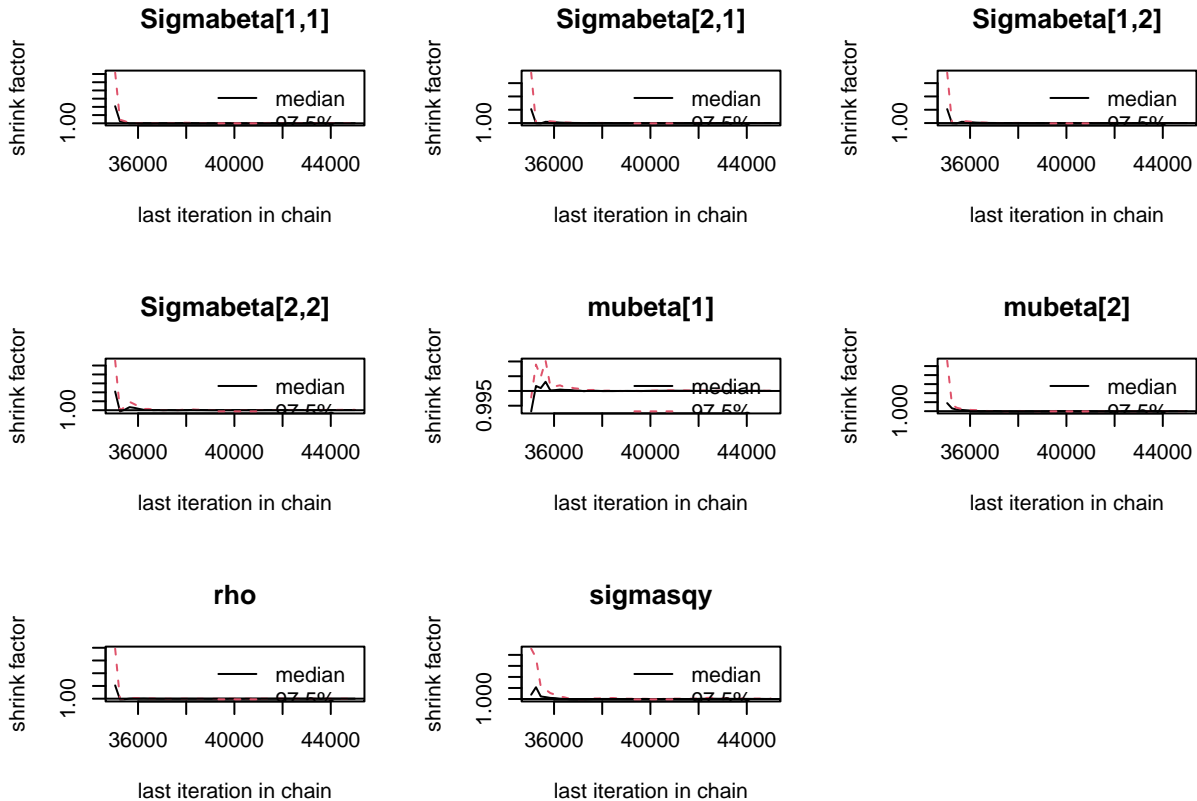
(ii) Display the coda summary of the results for the monitored parameters.

```
update(m1, 35000) # burn-in
x1 <- coda.samples(m1, c("mubeta", "Sigmabeta", "sigmasqy", "rho"), n.iter=10000)

gelman.diag(x1, autoburnin=FALSE, multivariate=FALSE)
```

```
## Potential scale reduction factors:
##
##           Point est. Upper C.I.
## Sigmabeta[1,1]           1           1
## Sigmabeta[2,1]           1           1
## Sigmabeta[1,2]           1           1
## Sigmabeta[2,2]           1           1
## mubeta[1]                 1           1
## mubeta[2]                 1           1
## rho                       1           1
## sigmasqy                  1           1
```

```
gelman.plot(x1, c("mubeta","Sigmabeta","sigmasqy", "rho"), autoburnin=FALSE)
```



```
effectiveSize(x1)
```

```
## Sigmabeta[1,1] Sigmabeta[2,1] Sigmabeta[1,2] Sigmabeta[2,2]      mubeta[1]
##      38198.64      38389.14      38389.14      37984.53      40726.68
##      mubeta[2]              rho      sigmasqy
##      40000.00      38371.68      22209.47
```

```
summary(x1[,c("mubeta[1]", "mubeta[2]", "Sigmabeta[1,1]", "Sigmabeta[1,2]",
              "Sigmabeta[2,2]", "sigmasqy", "rho")])$statistics
```

```
##              Mean              SD      Naive SE Time-series SE
## mubeta[1]      8.160194389 0.3655607441 1.827804e-03 1.811932e-03
## mubeta[2]      0.019886477 0.0056703000 2.835150e-05 2.835235e-05
## Sigmabeta[1,1] 5.337944385 1.2484454503 6.242227e-03 6.388376e-03
## Sigmabeta[1,2] -0.005834996 0.0136789745 6.839487e-05 6.982092e-05
## Sigmabeta[2,2] 0.001291679 0.0003046300 1.523150e-06 1.564426e-06
## sigmasqy      0.009867360 0.0009941911 4.970955e-06 6.674398e-06
## rho          -0.069189490 0.1551427443 7.757137e-04 7.920312e-04
```

```
summary(x1[,c("mubeta[1]", "mubeta[2]", "Sigmabeta[1,1]", "Sigmabeta[1,2]",
              "Sigmabeta[2,2]", "sigmasqy", "rho")])$quantiles
```


```
##              2.5%              25%              50%              75%              97.5%
## mubeta[1]      7.436360860 7.917841372 8.159036179 8.405882025 8.876009106
## mubeta[2]      0.008637544 0.016158242 0.019879969 0.023624875 0.031082424
## Sigmabeta[1,1] 3.418676524 4.452191897 5.165893634 6.025735764 8.269149271
## Sigmabeta[1,2] -0.034176613 -0.014162094 -0.005559062 0.002829731 0.020678222
## Sigmabeta[2,2] 0.000825890 0.001077522 0.001248754 0.001456135 0.002011983
```

```
## sigmasqy      0.008116119  0.009168586  0.009798569  0.010498616  0.011986597
## rho          -0.366604770 -0.176387227 -0.071420504  0.036303220  0.238140476
```

(iii) Give an approximate 95% central posterior interval for the correlation parameter  $\rho$ , and also produce a graph of its (estimated) posterior density.

An approximate 95% central posterior interval for the correlation parameter  $\rho$  is  $(-0.3666, 0.2381)$ .

```
densplot(x1[,c("rho")])
```



N = 10000 Bandwidth = 0.01975

(iv) Approximate the posterior probability that  $\rho < 0$ . Also, approximate the Bayes factor favoring  $\rho < 0$  versus  $\rho \geq 0$ . Describe the level of data evidence for  $\rho < 0$ .

```
# posterior probability for rho < 0
post.pro = mean(as.matrix(x1[,c("rho")]) < 0)
# Bayes factor
factor <- mean(as.matrix(x1[,c("rho")]) < 0) / (mean(as.matrix(x1[,c("rho")]) >= 0))

cat("The posterior probability for rho < 0 is:", post.pro, ",
    \n and Bayes factor favoring rho < 0 over rho >= 0 is:", factor,
    ", \n which means the data evidence is barely mentionable. \n")
```

```
## The posterior probability for rho < 0 is: 0.673 ,
##
## and Bayes factor favoring rho < 0 over rho >= 0 is: 2.058104 ,
## which means the data evidence is barely mentionable.
```

(v) The model implies that, over the 6 decades from 1950 to 2010, the median population change factor should be  $e^{6\mu_{\beta_2}}$ . Form an approximate 95% central posterior interval for this factor.

```
cat("An approximate 95% central posterior interval for the median population change factor is:
    \n (" , exp(6 * 0.0087), ",", exp(6 * 0.0311), "). \n")
```

```
## An approximate 95% central posterior interval for the median population change factor is:
##
```



```
## ( 1.053586 , 1.205145 ).
```

(vi) Use the `rjags` function `dic.samples` to compute the effective number of parameters (“penalty”) and Plummer’s DIC (“Penalized deviance”). Use at least 100,000 iterations.

```
dic.samples(m1, 200000)
```

```
## Mean deviance: -499.9
## penalty 81.72
## Penalized deviance: -418.2
```

(d) Now consider a different model with “univariate” hyperpriors for the model coefficients, which do not allow for a coefficient correlation parameter:

$$\beta_1^{(j)} | \mu_{\beta_1}, \sigma_{\beta_1} \sim \text{iid } N(\mu_{\beta_1}, \sigma_{\beta_1}^2)$$

$$\beta_2^{(j)} | \mu_{\beta_2}, \sigma_{\beta_2} \sim \text{iid } N(\mu_{\beta_2}, \sigma_{\beta_2}^2)$$


with hyperpriors

$$\mu_{\beta_1}, \mu_{\beta_2} \sim \text{iid } N(0, 1000^2)$$

$$\sigma_{\beta_1}, \sigma_{\beta_2} \sim \text{iid } U(0, 1000)$$

(i) Draw a complete DAG for this new model.

```
library(png)
img <- readPNG("dag.png")
plot(NA, xlim = c(0, 1.5), ylim = c(0, 1.1), type = "n", xlab = "", ylab = "", xaxt = "n",
     yaxt = "n", bty = "n", asp = 1)
rasterImage(img, 0, 0, 1.9, 1.1)
```



(ii) List an appropriate JAGS model. Make sure that there are nodes for  $\sigma_{\beta_1}$ ,  $\sigma_{\beta_2}$ , and  $\sigma_y^2$ .

```
d2 <- list(population = data[,-1],
           year = xs)
```

```

inits2 <- list(list(sigmatqyinv = 10, mubeta1 = 1000, mubeta2 = 1000,
                  sigmabeta1 = 1000, sigmabeta2 = 1000),
              list(sigmatqyinv = 0.001, mubeta1 = -1000, mubeta2 = 1000,
                  sigmabeta1 = 1000, sigmabeta2 = 1000),
              list(sigmatqyinv = 10, mubeta1 = 1000, mubeta2 = -1000,
                  sigmabeta1 = 0.001, sigmabeta2 = 0.001),
              list(sigmatqyinv = 0.001, mubeta1 = -1000, mubeta2 = -1000,
                  sigmabeta1 = 0.001, sigmabeta2 = 0.001))

library(rjags)

#print the model
cat (readLines('countrypop_2.bug'), sep= '\n')

## data {
##   dimY <- dim(population)
##   yearcent <- year - mean(year)
## }
##
## model {
##   for (j in 1:dimY[1]) {
##     for (i in 1:dimY[2]) {
##       population[j,i] ~ dnorm(beta[1,j] + beta[2,j]*yearcent[i], sigmasqyinv)
##     }
##     beta[1,j] ~ dnorm(mubeta1, sigmabeta1sqinv)
##     beta[2,j] ~ dnorm(mubeta2, sigmabeta2sqinv)
##   }
##   mubeta1 ~ dnorm(0, 0.000001)
##   mubeta2 ~ dnorm(0, 0.000001)
##   sigmabeta1 ~ dunif(0, 1000)
##   sigmabeta2 ~ dunif(0, 1000)
##   sigmasqyinv ~ dgamma(0.0001, 0.0001)
##
##   sigmabeta1sqinv <- 1/sigmabeta1^2
##   sigmabeta2sqinv <- 1/sigmabeta2^2
##   sigmasqy <- 1/sigmatqyinv
## }

m2 <- jags.model("countrypop_2.bug", d2, inits2, n.chains=4, n.adapt=3000)

## Compiling data graph
##   Resolving undeclared variables
##   Allocating nodes
##   Initializing
##   Reading data back into data table
## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 280
##   Unobserved stochastic nodes: 85
##   Total graph size: 952
##
## Initializing model

```

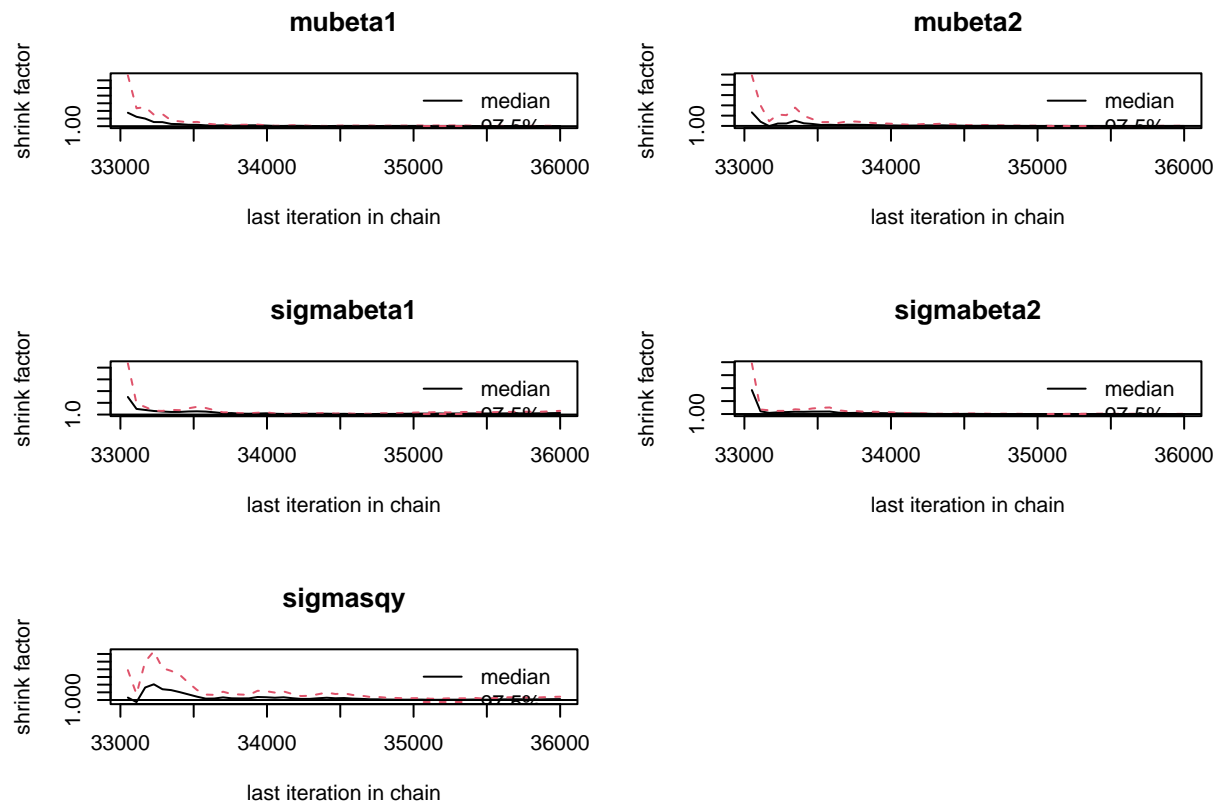
(iii) Display the coda summary of the results for the monitored parameters.

```
update(m2, 30000) # burn-in
x2 <- coda.samples(m2, c("mubeta1", "mubeta2", "sigmabeta1", "sigmabeta2", "sigmasqy"), n.iter=3000)
gelman.diag(x2, autoburnin=FALSE, multivariate=FALSE)
```

```
## Potential scale reduction factors:
```

```
##
##           Point est. Upper C.I.
## mubeta1           1.00      1.00
## mubeta2           1.00      1.00
## sigmabeta1        1.01      1.03
## sigmabeta2        1.00      1.00
## sigmasqy          1.00      1.00
```

```
gelman.plot(x2, autoburnin=FALSE)
```



```
effectiveSize(x2)
```

```
##      mubeta1      mubeta2 sigmabeta1 sigmabeta2      sigmasqy
## 13993.993 11554.684   4769.414   9494.022   6709.319
```

```
summary(x2)$statistics
```

```
##           Mean           SD      Naive SE Time-series SE
## mubeta1  8.155365306 0.3603122344 3.289186e-03 3.063226e-03
## mubeta2  0.019886437 0.0024989401 2.281210e-05 2.326968e-05
## sigmabeta1 2.281863860 0.2555192557 2.332561e-03 9.107543e-03
## sigmabeta2 0.015619229 0.0018442580 1.683570e-05 1.896227e-05
## sigmasqy  0.009898141 0.0009978783 9.109341e-06 1.223398e-05
```

```
summary(x2)$quantiles
```

```
##              2.5%      25%      50%      75%      97.5%
## mubeta1    7.450848172 7.916140833 8.155707468 8.39519273 8.86460378
## mubeta2     0.014993240 0.018235940 0.019866919 0.02153941 0.02482687
## sigmabeta1 1.851591931 2.100474872 2.257439450 2.43847124 2.84876203
## sigmabeta2 0.012493525 0.014316361 0.015454024 0.01670218 0.01974435
## sigmasqy   0.008131373 0.009193654 0.009835128 0.01051720 0.01201931
```

(iv) Form an approximate 95% central posterior interval for median change factor, and compare it with the previous results.

```
cat("An approximate 95% central posterior interval for the median change factor is:
    \n (", exp(6 * 0.0150), ",", exp(6 * 0.0248), "). \n")
```

```
## An approximate 95% central posterior interval for the median change factor is:
##
## ( 1.094174 , 1.160441 ).
```

Comparing with the previous result (1.053586 , 1.205145), the new one has similar but relatively more restrictive interval.

(v) Use the rjags function `dic.samples` to compute the effective number of parameters (“penalty”) and Plummer’s DIC (“Penalized deviance”). Use at least 100,000 iterations.

```
dic.samples(m2, 100000)
```

```
## Mean deviance: -499.8
## penalty 81.26
## Penalized deviance: -418.5
```

(vi) Compare the (Plummer’s) DIC values for this model and the previous one (-418.3), the conclusion is: These two models have very similar DICs, so it is hard to tell which model would be preferred by only using DICs.